

Laboratory Exercise 10

An Enhanced Processor

In Laboratory Exercise 1 we described a simple processor. In Part I of that exercise the processor itself was designed, and in Part II the processor was connected to an external counter and a memory unit. This exercise describes subsequent parts of the processor design. The numbering of figures and tables in this exercise are continued from those in Parts I and II of the preceding lab exercise.

In this exercise we will extend the capability of the processor so that the external counter is no longer needed, and so that the processor can perform read and write operations using memory or other devices. A schematic of the enhanced processor is given in Figure 12. In the figure registers $r0$ to $r6$ are the same as in Figure 1 of Lab 1, but register $r7$ has been changed to a counter. This counter is used to provide the addresses in the memory from which the processor's instructions are read; in the preceding lab exercise, a counter external to the processor was used for this purpose. We will refer to $r7$ as the processor's *program counter* (pc), because this terminology is common for real processors available in the industry. When the processor is reset, pc is set to address 0. At the start of each instruction (in time step T_0) the value of pc is used as an address to read an instruction from the memory. The instruction returned from the memory is stored into the IR register and the pc is automatically incremented to point to the next instruction.

The processor's control unit increments pc by using the $incr_pc$ signal, which is just an enable on this counter. It is also possible to load an arbitrary address into pc by having the processor execute an instruction in which the destination register is specified as pc . In this case the control unit uses $r7_{in}$ to perform a *load* of the counter. Thus, the processor can execute instructions at any address in the memory, as opposed to only being able to execute instructions that are stored at successive addresses. The current content of pc , which always has the address of the *next* instruction to be executed, can be copied into another register if needed by using a *mv* instruction.

The enhanced processor will have four new instructions, which are listed in Table 3. The *ld* (load) instruction *reads* data into register rX from the external memory address specified in register rY . Thus, the syntax $[rY]$ means that the content of register rY is used as an *external address*. The *st* (store) instruction *writes* the data contained in register rX into the memory address found in rY . The *and* instruction is similar to the *add* and *sub* instructions that were introduced in Lab 1. This instruction extends the adder/subtractor unit in the processor into an *arithmetic logic unit*. Besides performing addition and subtraction, it has the ability to generate a bit-wise logical AND ($\&$) of the destination register rX with the second operand $Op2$. As discussed in Lab 1, the operand $Op2$ can be either another register rY , or immediate data $\#D$.

The $b\{cond\}$ instruction is used to cause a processor *branch*, which means to change the program counter (pc) to the address of a specific instruction. The *cond* part of the branch instruction is optional and represents a *condition*. The instruction loads the constant $\#Label$ into pc only if the specified condition evaluates to true. An example of a condition is *eq*, which stands for *equal* (to zero). The instruction `beq #Label` will load the constant *Label* into pc if the last result produced by the arithmetic logic unit, which is stored in register G , was 0. The $b\{cond\}$ instruction is discussed in more detail in Part V of this exercise.

Operation	Function performed
<i>ld</i> $rX, [rY]$	$rX \leftarrow [rY]$
<i>st</i> $rX, [rY]$	$[rY] \leftarrow rX$
<i>and</i> $rX, Op2$	$rX \leftarrow rX \& Op2$
$b\{cond\}$ $\#Label$	if (<i>cond</i>), $pc \leftarrow \#Label$

Table 3: New instructions in the enhanced processor.

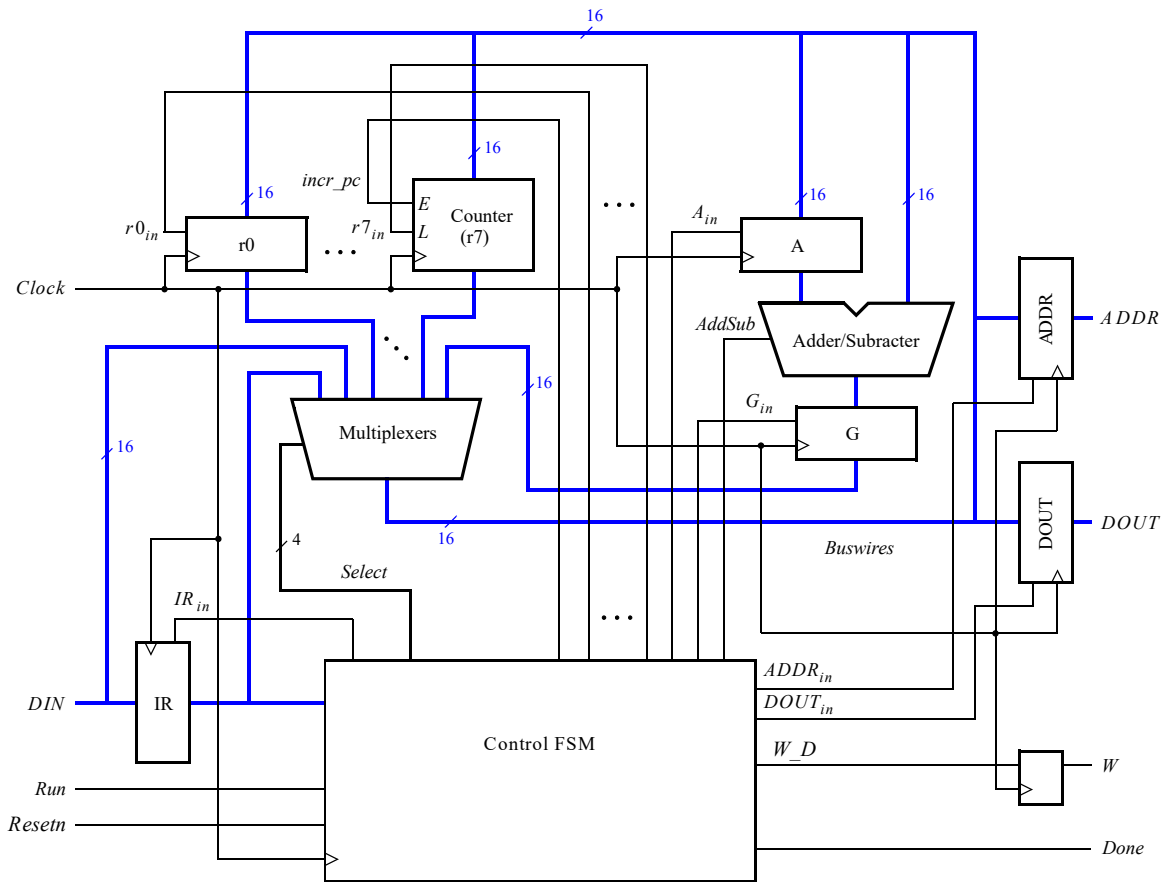


Figure 12: An enhanced version of the processor.

Recall from Lab 1 that instructions are encoded using a 16-bit format. For instructions that specify *Op2* as a register the encoding is `IIII0XXX000000YYY`, and if *Op2* is an immediate constant the format is `IIII1XXXDDDDDDDDDD`. You should use these same encodings for this exercise. Assume that `III` = 100 for the *ld* instruction, 101 for *st*, 110 for *and*, and 111 for *b{cond}*.

Figure 12 shows two registers in the processor that are used for data transfers. The *ADDR* register is used to send addresses to an external device, such as a memory unit, and the *DOUT* register is used by the processor to provide data that is to be stored outside of the processor. One use of the *ADDR* register is for reading, or *fetching*, instructions from memory; when the processor wants to fetch an instruction, the content of *pc* is transferred across the bus and loaded into *ADDR*. This address is provided to the memory.

In addition to fetching instructions, the processor can read data at any address by using the *ADDR* register. Both data and instructions are read into the processor on the *DIN* input port. The processor can write data for storage at an external address by placing this address into the *ADDR* register, placing the data to be stored into its *DOUT* register, and asserting the output of the *W* (*Write*) flip-flop to 1.

Connecting the Processor to External Devices

Figure 13 illustrates how the enhanced processor can be connected to memory and other devices. The memory unit in the figure is 16-bits wide and 256-words deep. A diagram of this memory is given in Figure 14. It supports both read and write operations and therefore has both address and data inputs, as well as a write-enable input. As

depicted in Figure 14, the memory has a clock input that is used to store the address, data, and write enable inputs into registers. This type of memory unit is called a *synchronous* static random access memory (SSRAM).

Figure 13 also includes a 9-bit output port (register) that can be used to store data from the processor. In the figure this output port is connected to a set of LEDs, like the ones available on the DE1-SoC board. To allow the processor to select either the memory unit or output port when performing a write operation, the circuit includes *address decoding*, which is done using NOR gates and AND gates. If the processor's upper address lines $A_{15}A_{14}A_{13}A_{12} = 0000$, then the memory unit can be written. Figure 13 shows n lower address lines connected from the processor to the memory; since the memory has 256 words, then $n = 8$ and the memory's *address* port is driven by the processor address lines $A_7 \dots A_0$. For addresses in which $A_{15}A_{14}A_{13}A_{12} = 0001$, the data written by the processor is loaded into the output port connected to LEDs in Figure 13.

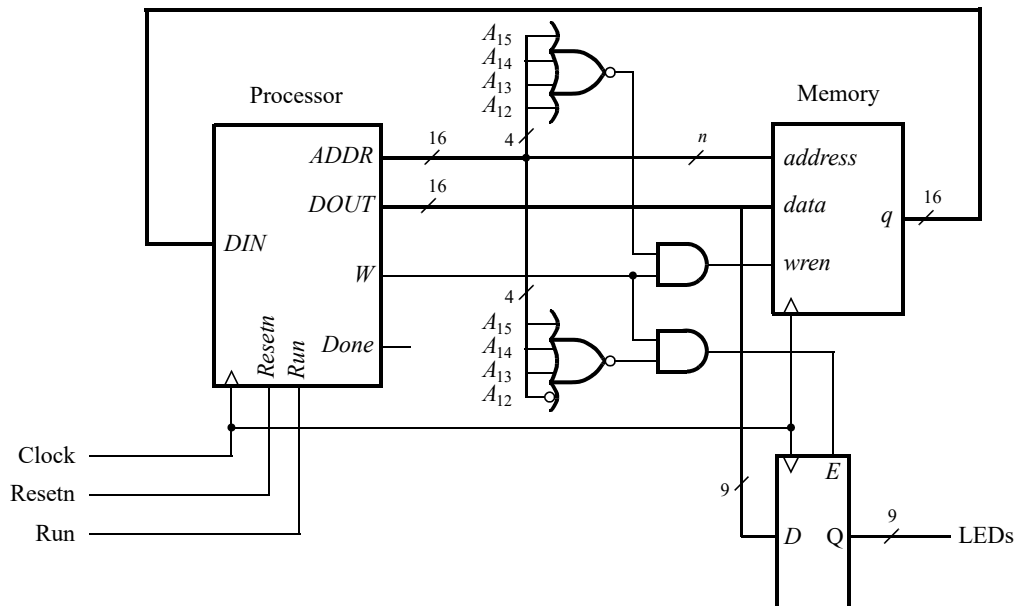


Figure 13: Connecting the enhanced processor to a memory unit and output register.




Figure 14: The synchronous SRAM unit.

Part III

Figure 15 gives VHDL code for a top-level file that you can use for this part of the exercise. The input and output ports for this entity are chosen so that it can be implemented on a DE1-SoC board. The VHDL code corresponds to the circuit in Figure 13, plus an additional input port that is connected to switches $SW_8 \dots SW_0$. This input port can be read by the processor at addresses in which $A_{15} \dots A_{12} = 0011$. (Switch SW_9 is not a part of the input port, because it is dedicated for use as the processor's *Run* input.) To support reading from both the SW input port and the memory unit, the top-level circuit includes a multiplexer that feeds the processor's *DIN* input. This multiplexer is described by using an *if-else* statement inside the *PROCESS* block in Figure 15.

The code in Figure 15 is provided with this exercise, along with a few other source-code files: *flipflop.v*, *inst_mem.v*, *inst_mem.mif*, and (part of) *proc.v*. The *inst_mem.v* source-code file was created by using the Quartus IP Catalog to instantiate a RAM:1-PORT memory unit. It has a 16-bit wide read/write data port and is 256-words deep, corresponding to Figure 14.

```
ENTITY part3 IS
    PORT ( KEY          : IN  STD_LOGIC_VECTOR(0 DOWNTO 0);
          SW            : IN  STD_LOGIC_VECTOR(9 DOWNTO 0);
          CLOCK_50      : IN  STD_LOGIC;
          LEDR          : OUT STD_LOGIC_VECTOR(9 DOWNTO 0) );
END part3;

ARCHITECTURE Behavior OF part3 IS
    ... declare components and signals
BEGIN
    U3: proc PORT MAP (DIN, KEY(0), CLOCK_50, Run, DOUT, ADDR, W);

    inst_mem_cs <= '1' WHEN (ADDR(15 DOWNTO 12) = "0000") ELSE '0';
    LED_reg_cs <= '1' WHEN (ADDR(15 DOWNTO 12) = "0001") ELSE '0';
    SW_cs <= '1' WHEN (ADDR(15 DOWNTO 12) = "0011") ELSE '0';

    W_mem <= inst_mem_cs AND W;
    U4: inst_mem PORT MAP (ADDR(7 DOWNTO 0), CLOCK_50, DOUT, W_mem, inst_mem_q);

    multiplexer: PROCESS (inst_mem_cs, SW_cs, inst_mem_q, SW_reg)
    BEGIN
        IF inst_mem_cs = '1' THEN
            DIN <= inst_mem_q;
        ELSIF SW_cs = '1' THEN
            DIN <= "0000000" & SW_reg;
        ELSE
            DIN <= (OTHERS => '-');
        END IF;
    END PROCESS;

    W_LED <= LED_reg_cs AND W;
    U5: regn GENERIC MAP (n => 9) PORT MAP (DOUT(8 DOWNTO 0), W_LED, CLOCK_50,
        LED_reg);
    LEDR(8 DOWNTO 0) <= LED_reg;
    LEDR(9) <= Run;
    High <= '1';
    U6: regn GENERIC MAP (n => 9) PORT MAP (SW(8 DOWNTO 0), High, CLOCK_50,
        SW_reg);
END Behavior;
```

Figure 15: VHDL code for the top-level file.

The VHDL code in the *proc.v* file implements register *r7* as a program counter, as discussed above, and includes a number of changes that are needed to support the new *ld*, *st*, *and*, and *b{cond}* instructions. In this part you are to augment this VHDL code to complete the implementation of the *ld* and *st* instructions, as well as the *and* instruction. You do not need to work on the *b{cond}* instruction for this part.

Perform the following:

1. Extend the code in *proc.v* so that the enhanced processor fully implements the *ld*, *st*, and *and* instructions. Test your VHDL code by using the ModelSim simulator. Sample setup files for ModelSim, including a testbench, are provided along with the other files for this exercise. The sample testbench first resets the processor system and then asserts the *Run* switch, *SW₉*, to 1. A sample program to test your processor is also provided, in a file called *inst_mem.mif*. This file represents the assembly-language program shown in Figure 16, which tests the *ld* and *st* instructions by reading the values of the SW switches and writing these values to the LEDs, in an endless loop. At the beginning of a simulation, ModelSim loads the contents of the file *inst_mem.mif* into the *inst_mem* memory unit, so that the program can be executed by the processor. Examine the signals inside your processor, as well as the external LEDR values, as the program executes within the ModelSim simulation.

An assembler software tool, called *sbasm.py*, can be used with your processor. The Assembler is written in Python and is available at <https://github.com/profbrown/sbasm.git>. To use this Assembler you have to first install Python (version 3) on your computer. The Assembler includes a README file that explains how to install and use it. The *sbasm.py* Assembler can generate machine code for all of the processor's instructions. The provided file *inst_mem.mif* was created by using *sbasm.py* to assemble the program in Figure 16. As the figure indicates, you can define symbolic constants in your code by using the *.define directive*, and you can use labels to refer to lines of code, such as *MAIN*. Comments are specified in the code by using *//*. The assembler ignores anything on a line following *//*.

```
.define LED_ADDRESS 0x1000
.define SW_ADDRESS 0x3000

// Read SW switches and display on LEDs
    mvt    r3, #LED_ADDRESS // point to LED port
    mvt    r4, #SW_ADDRESS  // point to SW port
MAIN:    ld     r0, [r4]      // read SW values
        st     r0, [r3]      // light up LEDs
        mv     pc, #MAIN
```

Figure 16: Assembly-language program that uses *ld* and *st* instructions.

An example result produced by using *ModelSim* for a correctly-designed circuit is given in Figure 17. It shows the execution of the first four instructions in Figure 16.

2. Once your simulation results are correct, use the Quartus Prime software to implement your VHDL code on a DE1-SoC board. A sample Quartus project file, *part3.qpf*, and Quartus settings file, *part3.qsf*, are provided with the exercise. Compile your code using the Quartus software, and download the resulting circuit into the DE1-SoC board. Toggle the SW switches and observe the LEDs to test your circuit.




Figure 17: Simulation results for the processor.

Part IV

In this part you are to create a new VHDL entity that represents an output port called *seg7*. It will allow your processor to write data to each of the six 7-segment displays on a DE1-SoC board. The *seg7* entity will include six write-only seven-bit registers, one for each display. Each register should directly drive the segment lights for one seven-segment display, so that the processor can write characters onto the displays.

Perform the following:

1. A top-level file is provided for this part called *part4.v*. The top-level entity has output ports for connecting to each of the 7-segment displays. Pin assignments for these ports, which are called *HEX0*[6:0], *HEX1*[6:0], ..., *HEX6*[6:0], are included in the Quartus settings file *part4.qsf*. For each display, segment 0 is on the top of the display, and then segments 1 to 5 are assigned in a clockwise fashion, with segment 6 being in the middle of the display.

The *part4.v* VHDL code includes address decoding for the new *seg7* module, so that processor addresses in which $A_{15}A_{14}A_{13}A_{12} = 0010$ select this module. The intent is that address 0×2000 should write to the register that controls display *HEX0*, 0×2001 should select the register for *HEX1*, and so on. For example, if your processor writes 0 to address 0×2000 , then the *seg7* module should turn off all of the segment-lights in the *HEX0* display; writing $0 \times 7f$ should turn on all of the lights in this display.

2. You are to complete the partially-written VHDL code in the file *seg7.v*, so that it contains the required six registers—one for each 7-segment display.
3. You can compile and test your VHDL code by using the ModelSim setup files that are provided for this part of the exercise. An *inst_mem.mif* file is also provided that corresponds to the assembly-language program shown in Figure 18. This program works as follows: it reads the *SW* switch port and lights up a seven-segment display corresponding to the value read on SW_{2-0} . For example, if $SW_{2-0} = 000$, then the digit **0** is shown on *HEX0*. If $SW_{2-0} = 001$, then the digit **1** is displayed on *HEX1*, and so on, up to the digit **5** which would be shown on *HEX5* if $SW_{2-0} = 101$.
4. Once your simulation results look correct, you should compile the provided Quartus project, and then download and test the circuit on a DE1-SoC board.

```

.define HEX_ADDRESS 0x2000
.define SW_ADDRESS 0x3000

// This program shows the digits 543210 on the HEX displays. Each digit has to
// be selected by using the SW switches.
        mv     r5, pc           // return address for subroutine
        mv     pc, #BLANK       // call subroutine to blank the HEX displays
MAIN:    mvt    r2, #HEX_ADDRESS // point to HEX port
        mv     r3, #DATA        // used to get 7-segment display pattern

        mvt    r4, #SW_ADDRESS  // point to SW port
        ld     r0, [r4]         // read switches
        and    r0, #0x7         // use only SW2-0
        add    r2, r0           // point to correct HEX display
        add    r3, r0           // point to correct 7-segment pattern

        ld     r0, [r3]         // load the 7-segment pattern
        st     r0, [r2]         // light up HEX display

        mv     pc, #MAIN

// subroutine BLANK
// This subroutine clears all of the HEX displays
// input: none
// returns: nothing
// changes: r0 and r1. Register r5 provides the return address
BLANK:   mv     r0, #0           // used for clearing
        mvt    r1, #HEX_ADDRESS // point to HEX displays
        st     r0, [r1]         // clear HEX0
        add    r1, #1
        st     r0, [r1]         // clear HEX1
        add    r1, #1
        st     r0, [r1]         // clear HEX2
        add    r1, #1
        st     r0, [r1]         // clear HEX3
        add    r1, #1
        st     r0, [r1]         // clear HEX4
        add    r1, #1
        st     r0, [r1]         // clear HEX5

        add    r5, #1
        mv     pc, r5           // return from subroutine

DATA:    .word 0b00111111       // '0'
        .word 0b00000110       // '1'
        .word 0b01011011       // '2'
        .word 0b01001111       // '3'
        .word 0b01100110       // '4'
        .word 0b01101101       // '5'

```

Figure 18: Assembly-language program that tests the seven-segment displays.

Part V

In this part you are to enhance your processor so that it implements the *b{cond}* instruction. The *conditions* supported by the processor are called *eq*, *ne*, *cc*, and *cs*, which means that the variations of the branch instruction are *b*, *beq*, *bne*, *bcc*, and *bcs*. The *b* instruction *always* branches. For example, `b #MAIN` loads the address MAIN into the program counter. The meanings of the conditional versions are explained below.

The instruction *beq* means *branch if equal* (to zero). This instruction performs a branch operation (i.e., loads the provided #LABEL into the program counter) if the most recent result of an instruction executed using the arithmetic logic unit (ALU), which is stored in register *G*, was 0. Similarly, *bne* means *branch if not equal* (to zero). It performs a branch only if the contents of *G* are not equal to 0. The instruction *bcc* stands for *branch if carry clear*. It branches if the last add/subtract operation did *not* produce a carry-out. The opposite branch condition, *bcs*, *branch if carry set*, performs a branch if the most recent add/sub generated a carry-out. To support the conditional branch instructions, you should create two *condition-code flags* in your processor. One flag, *z*, should have the value 1 when the ALU generates a result of zero; otherwise *z* should be 0. The other flag, *c*, should be 1 when the adder/subtractor in the ALU produces a carry-out; otherwise *c* should be 0. Thus, *c* should be 1 when an *add* instruction generates a carry-out, or when a *sub* operation requires a borrow for the most-significant bit. Your FSM controller should examine these flags in the appropriate clock cycles when executing the *b{cond}* instructions.

The branch instructions are encoded using the format `IIII1XXXXDDDDDDDDDD`, where DDDDDDDDDD is the branch address and XXX is the branch condition. Assume that conditions are encoded as *none* (always branch) = 000, *eq* = 001, *ne* = 010, *cc* = 011, and *cs* = 100.

Perform the following:

1. Enhance your processor so that it implements the condition-code flags *z* and *c*, and supports the *b{cond}* instruction. To help with testing and debugging of your processor, setup files for ModelSim are provided, including a testbench. It simulates your processor instantiated in the top-level file *part5.v*, which is the same as the one from Part IV. An example *inst_mem.mif* file is also provided, which corresponds to the program in Figure 19. This program is quite short, which makes it suitable for visual inspection of the waveforms produced by a ModelSim simulation. The program uses a sequence of instructions that test the various conditional branches. If the program reaches the line of code labelled DEAD, then at least one instruction has not worked properly. An example of ModelSim output for a correctly-working processor is given in Figure 20. It shows the processor executing instructions near the end of the code in Figure 19. The instruction that is completed at simulation time 1030 ns is `add r0, #1 (0x5001)`. As shown in the figure, this instruction causes the carry flag, *c*, to become 1. The next instruction loaded into *IR*, at time 1090 ns, is `bcs #0xC (0xF80C)`. Finally, the instruction loaded at 1170 ns is `b #0 (0xF000)`.

```
MAIN:    mv     r0, #2
LOOP:    sub     r0, #1           // subtract to test bne
        bne     #LOOP
        beq     #T1             // r0 == 0, test beq
T1:      mv     pc, #DEAD
        mvt     r0, #0xFF00
        add     r0, #0xFF        // r0 = 0xFFFF
        bcc     #T2             // carry = 0, test bcc
        mv     pc, #DEAD
T2:      add     r0, #1
        bcs     #T3             // carry = 1, test bcs
        mv     pc, #DEAD
T3:      b       #MAIN
DEAD:    mv     pc, #DEAD
```

Figure 19: Assembly-language program that uses various branches.


```

    mv    r6, #T2          // save address of next test
T2:    bne    #S1          // test bne; should take the branch!
    mv    pc, #FAIL
S1:    mv    r6, #C3       // save address of next test
C3:    beq    #FAIL        // test beq; should not take the branch
    add    r2, #2          // count the last two successful tests

    mv    r6, #T3         // save address of next test
    mv    r3, #ALLONES
    ld     r3, [r3]
    add    r3, #1          // set the c flag

T3:    bcc    #FAIL        // test bcc; should not take the branch!
    mv    r6, #C4         // save address of next test
C4:    bcs    #C5          // test bcs; should take the branch
    b      #FAIL          // Argh!
C5:    add    r2, #2        // count the last two successful tests

    mv    r6, #T4
    mv    r3, #0
    add    r3, r3          // clear carry flag

T4:    bcc    #S2          // test bcc; should take the branch!
    mv    pc, #FAIL
S2:    mv    r6, #C6       // save address of next test
C6:    bcs    #FAIL        // test bcs; should bot take the branch!
    add    r2, #2          // count the last two successes

// finally, test ld and st from/to memory
    mv    r6, #T5         // save address of next test
    mv    r4, #_LDTEST
    ld     r4, [r4]
    mv    r3, #0x1A5
    sub    r3, r4
T5:    bne    #FAIL        // should not take the branch!
    add    r2, #1          // increment success count

    mv    r6, #T6         // save address of next test
    mv    r3, #0x1A5
    mv    r4, #_STTEST
    st     r3, [r4]
    ld     r4, [r4]
    sub    r3, r4
T6:    bne    #FAIL        // should not take the branch!
    add    r2, #1          // increment success count

    mv    pc, #PASS
// Loop over the six HEX displays
FAIL:  mvt    r3, #LED_ADDRESS
    st     r6, [r3]        // show address of failed test on LEDs
    mv    r5, #_FAIL
    mv    pc, #PRINT
PASS:  mvt    r3, #LED_ADDRESS
    st     r2, [r3]        // show success count on LEDs
    mv    r5, #_PASS

```

Figure 21: Assembly-language program that tests various instructions. (Part b)

```

PRINT:  mvt    r4, #HEX_ADDRESS      // address of HEX0
        // We would normally use a loop counting down from 6
        // with bne to display the six letters. But in this
        // testing code we can't assume that bne even works!

        ld     r3, [r5]              // get letter
        st     r3, [r4]              // send to HEX display
        add    r5, #1                // ++increment character pointer
        add    r4, #1                // point to next HEX display
        ld     r3, [r5]              // get letter
        st     r3, [r4]              // send to HEX display
        add    r5, #1                // ++increment character pointer
        add    r4, #1                // point to next HEX display
        ld     r3, [r5]              // get letter
        st     r3, [r4]              // send to HEX display
        add    r5, #1                // ++increment character pointer
        add    r4, #1                // point to next HEX display
        ld     r3, [r5]              // get letter
        st     r3, [r4]              // send to HEX display
        add    r5, #1                // ++increment character pointer
        add    r4, #1                // point to next HEX display
        ld     r3, [r5]              // get letter
        st     r3, [r4]              // send to HEX display
        add    r5, #1                // ++increment character pointer
        add    r4, #1                // point to next HEX display

HERE:   mv     pc, #HERE

_PASS:  .word  0b0000000001011110    // d
        .word  0b0000000001111001    // E
        .word  0b0000000001101101    // S
        .word  0b0000000001101101    // S
        .word  0b0000000001110111    // A
        .word  0b0000000001110011    // P

_FAIL:  .word  0b0000000001011110    // d
        .word  0b0000000001111001    // E
        .word  0b0000000000111000    // L
        .word  0b0000000000110000    // I
        .word  0b0000000001110111    // A
        .word  0b0000000001110001    // F

ALLONES: .word  0xFFFF
_LDTEST: .word  0x1A5
_STTEST: .word  0x15A

```

Figure 21: Assembly-language program that tests various instructions. (Part c)

Part VI

Write an assembly-language program that displays a binary counter on the LED port. Initialize the counter to 0, and then increment the counter by one in an endless loop. You should be able to control the speed at which the counter is incremented by using nested delay loops, along with the SW switches. If the SW switches are set to

their maximum value, 0b11111111, then the delay loops should cause the counter to increment slowly enough so that each change in the counter can be visually observed on the LEDs. Lowering the value of the SW switches should make the counter increment more quickly up to some maximum speed.

You can assemble your program by using the *sbasm.py* assembler, and then run it on your processor system from Part V. To do this, use the output produced by *sbasm.py* to *overwrite* the file *inst_mem.mif* in the folder that holds your Quartus project for Part V. To make use of the new *inst_mem.mif* file you do not need to completely recompile your VHDL code from Part V. Instead, execute the Quartus command Processing > Update Memory Initialization File, to include the new *inst_mem.mif* file in your Quartus project. Next, select the Quartus command Processing > Start > Start Assembler to produce a new programming *bitstream* for your DE1-SoC board. Finally, use the Quartus Programmer to download the new bitstream onto your board. If the *Run* signal is asserted, your processor should execute the new program.

Part VII

Augment your assembly-language program from Part VI so that counter values are displayed on the seven-segment display port rather than on the LED port. You should display the counter values as decimal numbers from 0 to 65535. The speed of counting should be controllable using the SW switches in the same way as for Part VI. As part of your solution you may want to make use of the code shown in Figure 22. This code provides a subroutine that divides the number in register *r0* by 10, returning the quotient in *r1* and the remainder in *r0*. Dividing by 10 is a useful operation when performing binary-to-decimal conversion. The *DIV10* subroutine assumes that *r6* is set up to be used as a *stack pointer*. Register *r2* is saved on the stack at the beginning of the subroutine, and then restored before returning. This is done so that *r2* is not unnecessarily changed by the subroutine. A skeleton of the required code for this part is shown in Figure 23.

```
// subroutine DIV10
// This subroutine divides the number in r0 by 10
// The algorithm subtracts 10 from r0 until r0 < 10, and keeps count in r1
// This subroutine assumes that r6 can be used as a stack pointer
// input: r0
// returns: quotient Q in r1, remainder R in r0
DIV10:
    sub    r6, #1           // save registers that are modified
    st     r2, [r6]         // save on the stack

DLOOP:   mv     r1, #0       // init Q
    mv     r2, #9           // check if r0 is < 10 yet
    sub    r2, r0
    bcc    #RETDIV          // if so, then return

INC:     add    r1, #1       // but if not, then increment Q
    sub    r0, #10          // r0 -= 10
    b      #DLOOP           // continue loop

RETDIV:  ld     r2, [r6]     // restore from the stack
    add    r6, #1
    add    r5, #1           // adjust the return address
    mv     pc, r5           // return results
```

Figure 22: A subroutine that divides by 10

As described previously, assemble your program with *sbasm.py*, update your *MIF* file in the Quartus software, generate a new bitstream file by using the Quartus Assembler, and then download the new bitstream onto your DE1-SoC board to run your new program.

```

.define HEX_ADDRESS 0x2000
.define SW_ADDRESS 0x3000
.define STACK 256           // bottom of memory

// This program shows a decimal counter on the HEX displays
    mv    r6, #STACK        // stack pointer
    mv    r5, pc            // return address for subroutine
    mv    pc, #BLANK        // call subroutine to blank the HEX displays
MAIN:    mv    r0, #0        // initialize counter
LOOP:    mvt    r1, #HEX_ADDRESS // point to HEX port
    ...
    ... use a loop to extract and display each digit
    ...

// Delay loop for controlling the rate at which the HEX displays are updated
    ...
    ... read from SW switches, and use a nested delay loop
    ...

    add    r0, #1           // counter += 1
    bcc    #LOOP           // continue until counter overflows

    mv    r5, pc            // return address for subroutine
    mv    pc, #BLANK        // call subroutine to blank the HEX displays
    b      #MAIN

// subroutine DIV
    ...
    ... code not shown here
    ...
    add    r5, #1           // adjust the return address
    mv    pc, r5            // return results

// subroutine BLANK
    ...
    ... code not shown here
    ...
    add    r5, #1
    mv    pc, r5            // return from subroutine

DATA:    .word 0b00111111    // '0'
    ....

```

Figure 23: Skeleton code for displaying decimal digits.